Exercise 1 - Language Identification with sklearn

From Linear to Deep :-)

Deadlines

Deadline for Exercise 1 is **17.10.2021, 23:59 (Zurich Time)**.
Deadline for the peer review is **25.10.2021, 23:59 (Zurich Time).** You will find instructions for the peer review process at the end of this document.
Deadline for feedback to your peer reviewers is **29.10.2021, 23:59 (Zurich Time).**

Learning goals

This exercise consists of two parts: the first part aims at deepening your understanding of linear models. The second part will already target a simple kind of multi-layered network; the multilayer perceptron (MLP). Don't worry if you don't know anything about MLPs when you read this. We will cover all you need to know to do the second part of this exercise next week in class and in the tutorial session (hopefully). By completing this exercise you should …

- … understand linear models and use them for multiclass classification tasks.
- … be able to implement different machine learning models, including MLPs, in scikit-learn.
- … understand the role of hyper-parameters, regularisation and handling class imbalance.
- … perform an error analysis of machine learning models.

Please keep in mind that you can always consult and use the exercise forum if you get stuck (note that we have a separate forum for the exercises).

Deliverables

We request for you to hand in your solutions as ipynb files. That way your reviewers can view and execute your code without potential dependency problems and/or installing new packages or versions of packages.

Please hand in your solutions for part 1 and part 2 in separate notebooks. Name your notebooks exactly in the following fashion:

- ex01_lc.ipynb
- ex01_mlp.ipynb

The notebooks contain your well-documented, EXECUTED and EXECUTABLE code. You will also have to write a lab report. The lab report should contain a detailed description of the approaches you have used to solve this exercise. Please also include results. We highlight places where we expect a statement about an issue in your lab report.

We encourage you to solve the exercise on Google Colab and downloading the ipynb file when everything is run. Complimentary to the submission of the necessary ipynb files, you should also submit a lab report in a PDF format.

No matter if you decide to use colab or jupyter notebooks, please submit a zip-folder containing the following 3 files, named in exactly the following fashion:

- ex01_lc.ipynb
- ex01_mlp.ipynb
- ex01_labreport.pdf

We assume that the data is in the same folder as the scripts, e.g.

- ex01_ml4nlp1.zip
    - ex01_lc.ipynb
    - ex01_mlp.ipynb
    - Uniformly_sampled.tsv

Please note:

- The cells should be already run and the results should be visible to someone who will check your notebook without the need to rerun the code.
- Your peers need to be able to run your code. If it does not work, you will not be able to obtain the maximum number of points.
- DO NOT submit the data files!

Data

For both parts of this exercise, you will work with the same data. The goal is to classify the language of Tweets. This is an extension of the problem described in Goldberg, chapter 2. However, we will work with more languages than just six and the text segments we need to classify are much shorter. The material folder in the exercise section of OLAT contains the two files "train_dev_set.tsv" and "test_set.tsv". The files are also published under these two links:

- train_dev_set.tsv: https://docs.google.com/spreadsheets/d/e/2PACX-1vTOZ2rC82rhNsJduoyKYTsVeH6ukd7Bpxvxn_afOibn3R-eadZGXu82eCU9IRpl4CK_gefEGsYrA_oM/pub?gid=1863430984&single=true&output=tsv
- test_set.tsv: https://docs.google.com/spreadsheets/d/e/2PACX-1vT-KNR9nuYatLkSbzSRgpz6Ku1n4TN4w6kKmFLkA6QJHTfQzmX0puBsLF7PAAQJQAxUpgruDd_RRgK7/pub?gid=417546901&single=true&output=tsv

To make the start a little easier, you can go to this notebook in Google Colab which loads the two files using the public links. If you want to, you can just continue the exercise in your own copy of that notebook. If you choose to work locally, download the two files on your computer.

Inspect the data and see how it is distributed.

- What are the properties of the data? Describe the most important of distributional properties of this data set.

Use a 90/10 split for your train and development set[1]. Of course it is forbidden to peek into the test data ;-). You should only do this when you evaluate your model that performs best on the dev set.

Part 1 - Language identification with linear classification

Scikit-learn is a useful Python library for all kinds of machine learning tasks. In the following, you will train several models in sklearn to solve this task. The aim is to become acquainted with a few different classifiers, as well as with the basic functionality of sklearn.

1. Create a suitable pipeline in sklearn to preprocess the data. Think about extending the feature space. What other features could you use to determine the language? You're supposed to not only use bigrams for this task.
2. Train the following classifiers: SGDClassifier and Multinomial Naïve Bayes
3. In order to find the optimal hyperparameter settings for both classifiers, use sklearn's GridSearchCV. Especially with the SGDClassifier you are supposed to experiment with the following hyperparameters:
    a. Loss function
    b. Regularisation
    c. Early stopping
    d. Dealing with class imbalance
    e. Experiment with parameters of the Vectoriser

    Report the hyperparameter combination for your best performing model on the test set.

Also, compare the outputs of the best models for the two different classifiers. Which classifier scores higher on the test set? Do you have an idea, why this might be? What is the advantage of grid search cross validation? Use a confusion matrix to do your error analysis and summarise your answers in your report.

Part 2 - Your first Multilayer Perceptron (MLP)

Let's see if you can beat your best linear model you've trained with sklearn with a non-linear MLP.

1. Train an MLP classifier. You can also use GridSearchCV, but be aware that training an MLP model takes much more time. Play around with different layer sizes, activation functions, solvers, early stopping, vectoriser parameters and report your best hyperparameter combination. Important: use the same data splits as for Part 1.

If you need help on that, Raschka's [2015] chapter 2 provides an introduction to MLPs. The Google Machine Learning Crash Course also offers good material.

---

[1] Generally use the 90/10 split. When using SGD or MLP with early stopping, sklearn will create its own internal validation split. In these cases you are free to decide if you supply just the 90% training data or also the additional validation data.

Peer Review Instructions

First: Follow this link https://app.eduflow.com/join/TDY5CZ and register with the E-mail address you use for OLAT. Then you should be added to the course page automatically.
As soon as the deadline for handing in the exercise expires you will have time to review the submissions of your peers. You need to do **2 reviews** in order to get the maximum number of points for this exercise.
Here are some more rules:

- If you do not submit 2 reviews, the maximum number of points you can achieve is 0.75 (from a total of 1).
- Please use full sentences when giving feedback.
- Be critical, helpful, and fair!
- **All reviews are anonymous: Do not put your name into the python scripts, the lab report or the file names.**
- You must also give your reviewers feedback. The same criteria as above apply.
- Students that consistently provide very helpful feedback can be awarded a bonus in case they earned less than 6 points in total. Ways to obtain points are thus the following:
    - 5 exercises = 5 points
    - 1 presentation or research paper dissection = 1 points

Groups:

- You can create groups of two to solve the exercise together. This is strongly encouraged.
- When submitting the exercise, write a small post in the “Groups”-thread in the exercise forum on OLAT to notify the instructors about the group.
- As a group member you still have to review two submissions with your own eduflow account. However, you may work together in the group to write all 4 reviews.